

Inspection of Julia GC for multithreaded programs

Nazerke Turtayeva
Deept Mahendiratta

Contents

- Intro to Julia
- Syntax in short
- Motivation
- GC in Julia
- Tracing tools
- Julia GC tracing results
- Project challenges
- References

Intro to Julia

- What:** High-level, high-performance, dynamic programming language - aimed to employ the best of both worlds
- Who:** GP, but mainly for scientific community - numerical analysis
~10,000 companies, 1500 universities[1]
- When:** 2012, MIT
- Features:** Multiple dispatch, dynamic typing, parallel and distributed computing, no need for a vectorized code, package management
- Compiler:** Written in Julia, C/C++, JIT



vs



vs



Syntax in short

Q: What makes Julia to perform close to C?

A: (1) Julia features + (2) Efficient user code

Features:


- Dynamic type inferring

```
d = Int[1,2,3,4,5] vs d = [1,2,3,4,5]
```

- Multiple dispatch (don't overuse)
- Using function profiling

User:

- Type stability to infer types beforehand correctly
NO `d = 1 ... d = 2.3 ... d = [1,2,3,4,5]`
- Common sense: no global vars, static arrays, @inbounds, etc[2]
- Identifying bottlenecks



Types hierarchy in Julia [2]

```
function test1(x::Int)
    println("$x is an Int")
end

function test1(x::Float64)
    println("$x is a Float64")
end

function test1(x)
    println("$x is neither an Int
    nor a Float64")
end
```

Multiple dispatch example [2]

Motivation

Source of issues:

- 1)
 - > Non parallel, mark and sweep GC
 - > multi threading
 - > more data to collect during “stop the world”(stw)
 - > higher “stw” latency
 - > lower average performance
- 2)
 - > Some non-allocational threads might not never yield to a safe(precise) point
 - > Hence, allocating threads might stall
 - > lower average performance

Motivation

Solutions:

- 1) Inspecting the GC to identify which specific applications are more vulnerable to the “stw” latencies
- 2) Identifying which parts of the GC collections and sweeping relatable for this

GC in Julia

- Julia's garbage collector algorithm is called *mark and sweep*.
- It is precise, generational, non-compacting, non-parallel.
- The mark phase: Where all objects that are not garbage are found and marked so
- The sweep phase: where all unmarked objects are cleaned
- Starts earlier than memory overflow
- The top Julia GC source file is GC.c

Triggering manual GC

- When the garbage collector detects that the program can no longer access the object, then it will run the finalizer, and then collect (free) the object.
- Note that the garbage collector can still access the object, even though the program cannot.

Example of Manual GC

```
julia> r = Ref{0}  
Base.RefValue{Int64}(0)
```

```
julia> finalizer(r) do r  
    println(r)  
end  
Base.RefValue{Int64}(0)
```

```
julia> r = nothing
```

```
julia> GC.gc()  
Base.RefValue{Int64}(0)
```

DTrace and BPFTrace

- **DTrace** (Solaris originally) and **bpftrace** (Linux) are tools that enable **lightweight instrumentation** of processes.
- You can turn the instrumentation on and off while the process is running, and with instrumentation off the overhead is minimal.

Tracing interfaces

	kernel	userland
static	tracepoints	USDT* probes
dynamic	kprobes	uprobes

[5]

- We will focus on userland for now, since that is the most useful feature for understanding applications, like Julia
- USDT - **U**ser-level **S**tatically **D**efined **T**racing

DTrace and BPFTrace

- Using nm we can list all of the exported runtime functions of Julia.

```
nm -D /usr/lib/julia/libjulia-internal.so | grep jl_gc_
```

```
0000000000015215 T jl_gc_add_finalizer
0000000000015221 T jl_gc_add_finalizer_th
000000000001522d T jl_gc_add_ptr_finalizer
0000000000015239 T jl_gc_alloc
0000000000015245 T jl_gc_alloc_0w
0000000000015251 T jl_gc_alloc_1w
000000000001525d T jl_gc_alloc_2w
0000000000015269 T jl_gc_alloc_3w
0000000000015275 T jl_gc_allocobj
0000000000015281 T jl_gc_alloc_typed
000000000001528d T jl_gc_big_alloc
```

DTrace and BPFTrace

- An alternative way to find all valid uprobe is to use bpftrace:

```
sudo bpftrace -l 'uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_*
```

```
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_add_finalizer  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_add_finalizer_th  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_add_ptr_finalizer  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_0w  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_1w  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_2w  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_3w  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_page  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_alloc_typed  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_allocobj  
uprobe:/usr/lib/julia/libjulia-internal.so:jl_gc_big_alloc
```

Using uprobe[5]

- The function allocator is going to run forever and allocate an array of size N bytes.

```
function allocator(range)  
    while true  
        N = rand(range)  
        buf = Array{UInt8}(undef, N)  
    end  
end
```

- Running it: **julia -L allocator.jl -e "allocator(64:128)"**

Results: uprobe

```
ubuntu@euca-10-1-4-194:~/julia$ sudo bpftrace -e  
"uprobe:usr/lib/libjulia-internal.so:jl_gc_alloc { @[pid] = hist(arg1); }"  
Attaching 1 probe...  
^C
```

```
@[19075]:  
[64, 128)          100857 | @@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@ |  
[128, 256)         182855 | @@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@ |
```

- Memory allocation distribution by process ID
- This triggers all the processes using the library, but it is possible to trigger only by PID with `-p PID`

Using USDT probes[6]

- Enabling support:
Using `WITH_DTRACE=1` in `Make.user` file
- Available probes, `uprobes.d`:
 - `julia:gc__begin`: GC begins running on one thread and triggers stop-the-world.
 - `julia:gc__stop_the_world`: All threads have reached a safepoint and GC runs.
 - `julia:gc__mark__begin`: Beginning the mark phase
 - `julia:gc__mark__end(scanned_bytes, perm_scanned)`: Mark phase ended
 - `julia:gc__sweep__begin(full)`: Starting sweep
 - `julia:gc__sweep__end()`: Sweep phase finished
 - `julia:gc__end`: GC is finished, other threads continue work
 - `julia:gc__finalizer`: Initial GC thread has finished running finalizers

Using USDT probes

```
using Base.Threads
```

```
fib(x) = x <= 1 ? 1 : fib(x-1) + fib(x-2)
```

```
beaver = @spawn begin
```

```
    while true
```

```
        fib(30)
```

```
        GC.safepoint()
```

```
    end
```

```
end
```

```
allocator = @spawn begin
```

```
    while true
```

```
        zeros(1024)
```

```
    end
```

```
end
```

```
wait(allocator)
```

- Run this using `julia -t 2` (This means we are using 2 threads)
- `@spawn` allocates a Task to an available thread and executes

Results: USDT probe for “stop the world”

```
ubuntu@euca-10-1-4-194:~/julia$ sudo bpftrace --usdt-file-activation
contrib/bpftrace/gc stop_the_world_latency.bt
Attaching 4 probes...
Tracing Julia GC Stop-The-World Latency... Hit Ctrl-C to end.
^C
```

```
@usecs[16514]:  
[512, 1K)          1 |  
[1K, 2K)         105 | @@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@  
[2K, 4K)          65 | @@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@
```

- Latency distribution of the stop-the-world phase in the executed Julia process.

Contents of `gc_stop_the_world_latency.bt`

```
#!/usr/bin/env bpftrace
```

```
BEGIN
```

```
{  
    printf("Tracing Julia GC Stop-The-World Latency... Hit Ctrl-C to end.\n");  
}
```

```
usdt:usr/lib/libjulia-internal.so:julia:gc__begin
```

```
{  
    @start[pid] = nsecs;  
}
```

```
usdt:usr/lib/libjulia-internal.so:julia:gc__stop_the_world
```

```
/@start[pid]/  
{  
    @usecs[pid] = hist((nsecs - @start[pid]) / 1000);  
    delete(@start[pid]);  
}
```

```
END
```

```
{  
    clear(@start);  
}
```

Using USDT probes

```
using Base.Threads
```

```
fib(x) = x <= 1 ? 1 : fib(x-1) + fib(x-2)
```

```
beaver = @spawn begin
    while true
        fib(30)
        GC.safepoint()
    end
end
```

```
allocator = @spawn begin
    while true
        zeros(1024)
    end
end
```

```
wait(allocator)
```

- `GC.safepoint()` should be called from a non-allocating thread to make sure that all threads are checked for a precise state.
- It triggers segmentation fault by loading from a protected memory space under the hood and forces GC to execute.
- Future Julia versions expected to include built-in function entry `GC.safepoints()`

Project challenges

Never use M1 chip Mac for GC tracing! :)

Using linux servers with sudo access is a prerequisite

References

- [1] [https://en.wikipedia.org/wiki/Julia_\(programming_language\)#History](https://en.wikipedia.org/wiki/Julia_(programming_language)#History)
- [2] <https://techtok.com/code-optimisation-in-julia/>
- [3] <https://juliacomputing.com/>
- [4] <https://github.com/JuliaLang/julia>
- [5] <https://vchuravy.dev/notes/2021/08/bpftrace/>
- [6] <https://docs.julialang.org/en/v1.8-dev/devdocs/probes/#Available-probes>

UC SANTA BARBARA